

Algorithm Theory, Tutorial 7

Johannes Kalmbach

University of Freiburg

johannes.kalmbach@gmail.com

February 2020



General Hints

- Contact tutor (johannes.kalmbach@gmail.com) for questions concerning corrections etc.
- Contact forum (daphne.informatik.uni-freiburg.de) for everything else
- Suggestion: Submit in groups of two (better for understandable algorithms)
- Submit readable solutions (LaTeX as pdf, CLEAN handwriting (+ good scan if necessary))
- Spend enough time on exercise sheets and writeup (you and I have to understand your submission).

Exercise 1, Load Balancing

- Given m machines
- Given n jobs, each job i has processing time t_i
- Assign each job to a machine to minimize **makespan**
- Makespan: largest total processing time of any machine.
- *modified greedy* algorithm: Go through jobs by decreasing length, assign to machine with smallest current load ¹
- Modified Greedy has approximation ratio of $3/2$.
- In this exercise, we want to prove that the algorithm has an even better approximation ratio.

¹If there are machines with the same load, Greedy chooses the one with the smallest ID.

- Assume we have n jobs with lengths $t_1 \geq t_2 \geq \dots \geq t_n$.
- let i be a machine with load $T \leftarrow \text{makespan of greedy}$
- let \hat{n} be the last job that is scheduled on machine i .
- Shortly argue why it is sufficient to ignore jobs $\hat{n} + 1, \dots, n$ and instead prove the desired ratio between greedy and optimal for jobs $1, \dots, \hat{n}$.
-
- After scheduling job \hat{n} greedy already has makespan T (equal to the final makespan with all jobs).
- So ignoring the jobs after \hat{n} doesn't make the task easier for greedy.
- The optimal solution however might be better if we leave out these jobs.
- Thus the approximation ratio on the modified problem is \geq than on the original problem

Exercise 2b

- Assume we have n jobs with lengths $t_1 \geq t_2 \geq \dots \geq t_n$.
- let i be a machine with load T
- let \hat{n} be the last job that is scheduled on machine i .
- Show that if an optimal solution for jobs $1, \dots, \hat{n}$ assigns at most two jobs to each machine, the algorithm computes an optimal solution.

Hint: Think of a “canonical” way to assign at most two jobs to each machine and show that $T \leq T_{\text{canonical}} \leq T_{\text{opt}}$.

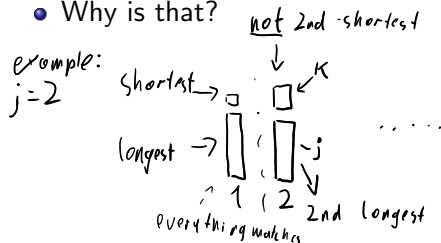
The canonical solution

- We know that there are $\leq 2m$ jobs.
- Fill up with empty jobs, s.t. there are exactly $2m$ jobs (does not change the problem, optimal solution with max 2 jobs/machine still exists).
- Put the longest and the shortest job on machine 1.
- Put the second longest and second shortest job on machine 2.
- Put the i -th and the $2m - i + 1$ -th job on machine i
- Claim: This is an optimal solution



Proof: Canonical is optimal

- Proof by exchange argument: Transform optimal into canonical solution without making it worse.
- Assume, j is the first index s.t. job j is not paired with job $2m - j + 1$ in the optimal solution.
- Job j is then paired with another job k and job $2m - j + 1$ is paired with another job j' .
- We can swap jobs $2m - j + 1$ and job k without increasing the makespan. Then we still have an optimal solution but one more index matches with the canonical solution.
- Why is that?



$2m-j+1$
 \square

$\square j'$

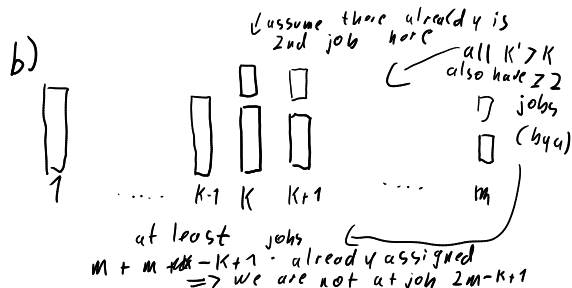
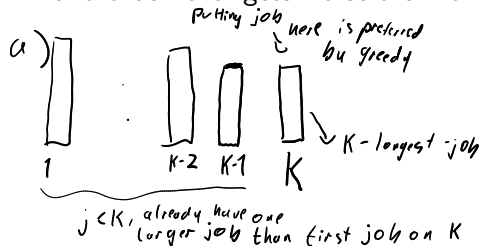
$$\begin{aligned} t_j + t_k &\geq t_{j'} + t_{2m-j+1} \\ t_j &\geq t_{j'} \\ t_k &\geq t_{2m-j+1} \end{aligned}$$

After swap:

$$\begin{aligned} t_j + t_{2m-j+1} &\leq t_j + t_k \\ t_{j'} + t_k &\leq t_{j'} + t_{2m-j+1} \end{aligned} \Rightarrow \text{swap is OK}$$

Greedy is optimal

- Our canonical solution is optimal, thus we have to show that $T_{\text{Greedy}} \leq T_{\text{Canonical}}$
- Observation: for the first m jobs, greedy and canonical do the same
- Claim: While performing greedy, if machine k has only one job, then all machines $j < k$ also have only one job. Why? a) $\hookleftarrow K\text{-shortest-job}$
- Thus: When greedy assigns job $2m - k + 1$, then machine k has only one job. Why? b)
- Thus, for each job greedy can always perform the same choice as the canonical solution, and thus never gets worse than it.



- Greedy can schedule each job j with its canonical partner \Rightarrow greedy never exceeds Makespan of Canonical = Optimal solution

Exercise 1c

- Assume we have n jobs with lengths $t_1 \geq t_2 \geq \dots \geq t_n$.
- let i be a machine with load T
- let \hat{n} be the last job that is scheduled on machine i .
- Show that therefore, either $t_{\hat{n}} \leq T_{opt}/3$ or the greedy algorithm computes an optimal solution.
- If there is an optimal solution with ≤ 2 jobs per machine, then greedy is optimal.
- Else, there is a machine in the optimal solution with at least 3 jobs.
- \hat{n} is the shortest job, and thus each of those jobs is $\geq t_{\hat{n}}$
- Thus there is machine that has a load of at least $3 \cdot t_{\hat{n}} \leq T_{opt}$, and thus $t_{\hat{n}} \leq T_{opt}/3$

\downarrow job/machine \nwarrow min. job duration

Exercise 1d

- Assume we have n jobs with lengths $t_1 \geq t_2 \geq \dots \geq t_n$.
- let i be a machine with load T
- let \hat{n} be the last job that is scheduled on machine i .
- Conclude that the algorithm has an approximation ratio of at most $4/3$.
- We only have to look at the case, where the optimal solution has at least three jobs for at least one machine, otherwise we are optimal.
- Greedy assigns job \hat{n} to the machine with minimal load. *↖ Last job*
- Minimal load before adding $\hat{n} \leq \text{Average load before adding} \leq T_{opt}$
- Since $t_{\hat{n}} \leq T_{opt}/3$ we have

$$T_{greedy} \leq 4/3 T_{opt} = \underbrace{T_{opt}}_{\substack{\text{upper bound} \\ \text{load of machine } i \\ \text{before adding } \hat{n}}} + \underbrace{\frac{T_{opt}}{3}}_{\substack{\text{upper bound} \\ \text{for } t_{\hat{n}}}}$$

And this machine has maximum load after forwards

- Show that the $4/3$ bound is tight, i.e., there is a sequence of instances for which the ratio between Greedy and OPT converges to $4/3$.

Hint: Consider $2m + 1$ jobs for m machines, three jobs with processing time m and two jobs with processing times $m + 1, m + 2, \dots, 2m - 1$ each.

- 1st job | 2nd job | (so far each machine has load $3m-1$, still have to add 1 job of length m to any of the machines \Rightarrow greedy value is $4m-1$)

m machines

$$\begin{array}{l}
 2m-1 + m \\
 2m-1 + m \\
 2m-2 + m+1 \\
 2m-2 + m+1 \\
 \dots \\
 1.5m \quad 1.5m-1 \\
 1.5m \quad 1.5m-1
 \end{array}$$

Optimal:

$$\begin{array}{l}
 m + m + m \\
 2m-1 + m+1 \\
 2m-1 + m+1 \\
 2m-2 + m+2 \\
 \dots
 \end{array}$$

\Rightarrow Approx Ratio $\frac{4m-1}{3m} \xrightarrow{m \rightarrow \infty} \frac{4}{3}$

\Rightarrow Optimal Solution is $3m$

Consider the following variation of the knapsack problem: Given items $1, \dots, n$ where each item i has a positive integer *weight* $w_i \in \mathbb{N}$ and a positive *value* $v_i > 0$ and **two** knapsacks of capacities W_1 and W_2 , we want to pack the items into the knapsacks such that

- for $j \in \{1, 2\}$, the *total weight* of the items in knapsack j is at most W_j .
- The *total value* of the items that are packed in either knapsack is maximized.

- Prove that this problem is not equivalent to the standard knapsack problem with one knapsack of capacity $W_1 + W_2$ by showing that the total value that can be packed into one knapsack of capacity $W_1 + W_2$ can be *arbitrarily* larger than the total value that can be packed into two knapsacks of capacities W_1 and W_2 .

$$W_1 = 1$$

$$W_2 = 1$$

\nearrow 2 knapsacks solution value = 0
 \searrow $W_1 + W_2$ solution value = x

1 item with weight 2, value x
 \Downarrow
 arbitrarily big

- Assume that $W_1 \geq W_2$. A simple strategy would be to first compute an optimal solution for a knapsack of capacity W_1 and afterwards, with the remaining elements, an optimal solution for a knapsack of capacity W_2 . Show that this algorithm always computes at least a 2-approximation for the problem.
- Let OPT_{W_1} be the content of W_1 in the optimal solution of the overall problem OPT_{W_2} similarly.
- let $OPT_{max} := \max(OPT_{W_1}, OPT_{W_2})$
- We know that $OPT_{max} \geq 0.5(OPT_{W_1} + OPT_{W_2})$ \leftarrow *max \geq average*
- Also we know that OPT_{max} fits into W_1 (because it is the bigger knapsack). *(feasible for knapsack on W_1)*
- Let $OPTSINGLE_{W_1}$ be an optimal solution on the "standard knapsack" on W_1 .
- We know that $OPTSINGLE_{W_1} \geq OPT_{max} \geq 0.5(OPT_{W_1} + OPT_{W_2})$
- So even calculating the traditional knapsack on W_1 and ignoring W_2 is a 2-approximation, so this also holds if we put some of the rest into W_2

Exercise 3

- Show that taking all nodes is a 2-approximation algorithm for the vertex cover problem in regular graphs (graphs where all nodes have the same degree)
- A r -regular graph has $\frac{n \cdot r}{2}$ ^{edges} ~~nodes~~ _(additional)
- each of the n nodes covers at most r ^{edges}
- so to cover all edges, we need at least $\frac{n}{2}$ nodes (every vertex cover on regular graphs has size $\geq \frac{n}{2}$)
- Taking all n nodes is a 2-approximation

